Internal Assessment Task -III - **Implementation of current trends in AI**

"Credit Default Prediction"

by

16010123247: Rahul Suhas Kadam

For the subject

Artificial Intelligence

Department of Computer Engineering
Academic Year 2025-26

Credit Default Prediction using Neural Networks

Problem Statement No. - 28

Google Colab Notebook Link –

https://colab.research.google.com/drive/17eRavA2tm1JkNaUmWh2w002DaoCB22AS?usp=sharing

Introduction, Abstract of problem statement

The problem statement - "To Mathematically prescribed credit limits for new client profiles to minimize default risk."

- ***Credit risk assessment is one of the most critical tasks in the financial industry. Banks and financial institutions must determine appropriate credit limits for new clients while minimizing the probability of default. Traditionally, this process relied heavily on rule-based systems and human judgment, which can be inconsistent and prone to bias.***
- ***In this project, we address the problem of mathematically prescribing credit limits for new client profiles using machine learning techniques. Specifically, we design and train a neural network model that learns patterns from historical credit card client data and predicts whether a client is likely to default. Based on this prediction, the system can assist in assigning safer credit limits.***
- ***The solution leverages the Default of Credit Card Clients dataset and applies preprocessing, feature engineering, and deep learning techniques to build an accurate predictive model. The final model achieves an accuracy greater than 75%, making it a reliable tool for decision support in credit risk management.***

Problem Definition and objectives of work

The core problem is to predict the likelihood of default for a credit card client and use that prediction to guide credit limit decisions.

Objectives:

- ***To analyze client financial and demographic data.***
- ***To preprocess and clean the dataset for model training.***
- ***To build a neural network capable of classifying default risk.***
- ***To achieve model accuracy greater than 75%.***
- ***To visualize the trained neural network architecture.***
- ***To evaluate the model using various performance metrics.***
- ***To test the model on custom inputs for real-world applicability.***

Datasets used, size, preprocessing if any

Dataset Overview -

- ***Dataset: Default of Credit Card Clients (UCI Repository)***
- ***Total Records: 30,000***
- ***Features: 23 input features + 1 target variable***
- ***Target Variable: default.payment.next.month (0 = No Default, 1 = Default)***

Column Description -

Since the dataset had unclear column names, they were properly defined as:

- ***LIMIT_BAL: Credit limit***
- ***SEX: Gender***
- ***EDUCATION: Education level***
- ***MARRIAGE: Marital status***
- ***AGE: Age***
- ***PAY_0 to PAY_6: Repayment status (past months)***
- ***BILL_AMT1 to BILL_AMT6: Bill statements***
- ***PAY_AMT1 to PAY_AMT6: Previous payments***

Preprocessing Steps

1. ***Handling Missing Values***

    - ***Dataset was clean, no missing values required handling.***
2. ***Feature Scaling***
    - ***Applied StandardScaler to normalize numerical features.***
    - ***Ensures faster convergence during training.***
3. ***Encoding Categorical Variables***
    - ***Converted categorical columns into numerical format.***
4. ***Train-Test Split***
    - ***80% training data***
    - ***20% testing data***
5. ***Balancing***
    - ***Checked class imbalance (default vs non-default).***

Architectural diagram if applicable/possible

A correlation heatmap was generated to understand relationships between variables.

- ***Strong correlations observed among bill amounts (BILL_AMT1–6).***
- ***Payment delays (PAY_0, etc.) showed positive correlation with default.***
- ***Higher credit limit (LIMIT_BAL) generally showed negative correlation with default.***

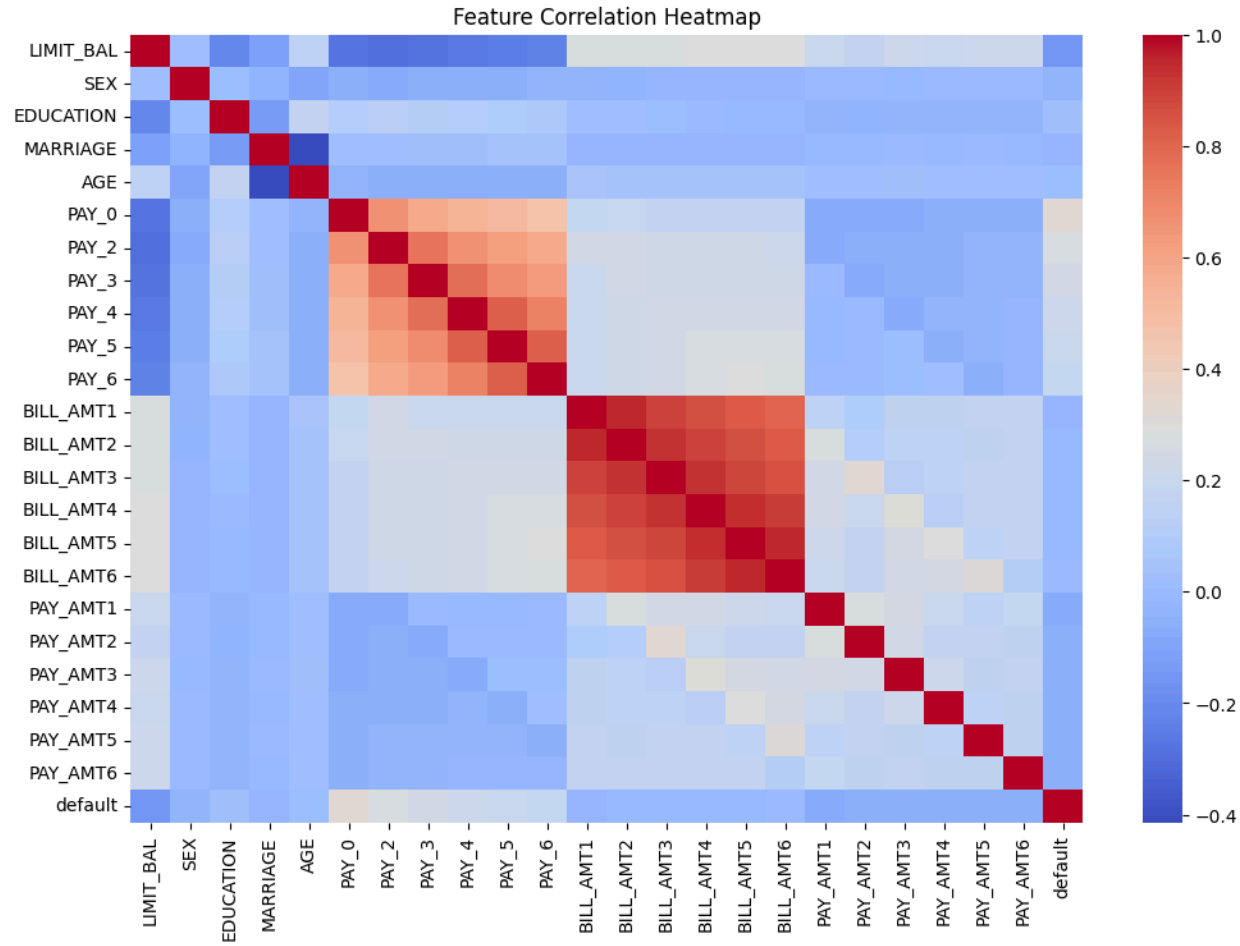

Clients with frequent delayed payments and high outstanding bills are more likely to default.

The neural network architecture as made by the netron neural network visualizer consists of:

- ***Input Layer (28 features)***
- ***Dense Layer (128 neurons)***
- ***Batch Normalization***
- ***Dropout - To prevent overfitting of the model.***
- ***Dense Layer (64 neurons)***
- ***Batch Normalization***
- ***Dropout***
- ***Dense Layer (32 neurons)***
- ***Output Layer (1 neuron)***

Below is the ScreenShot of the output given by Netron -

- 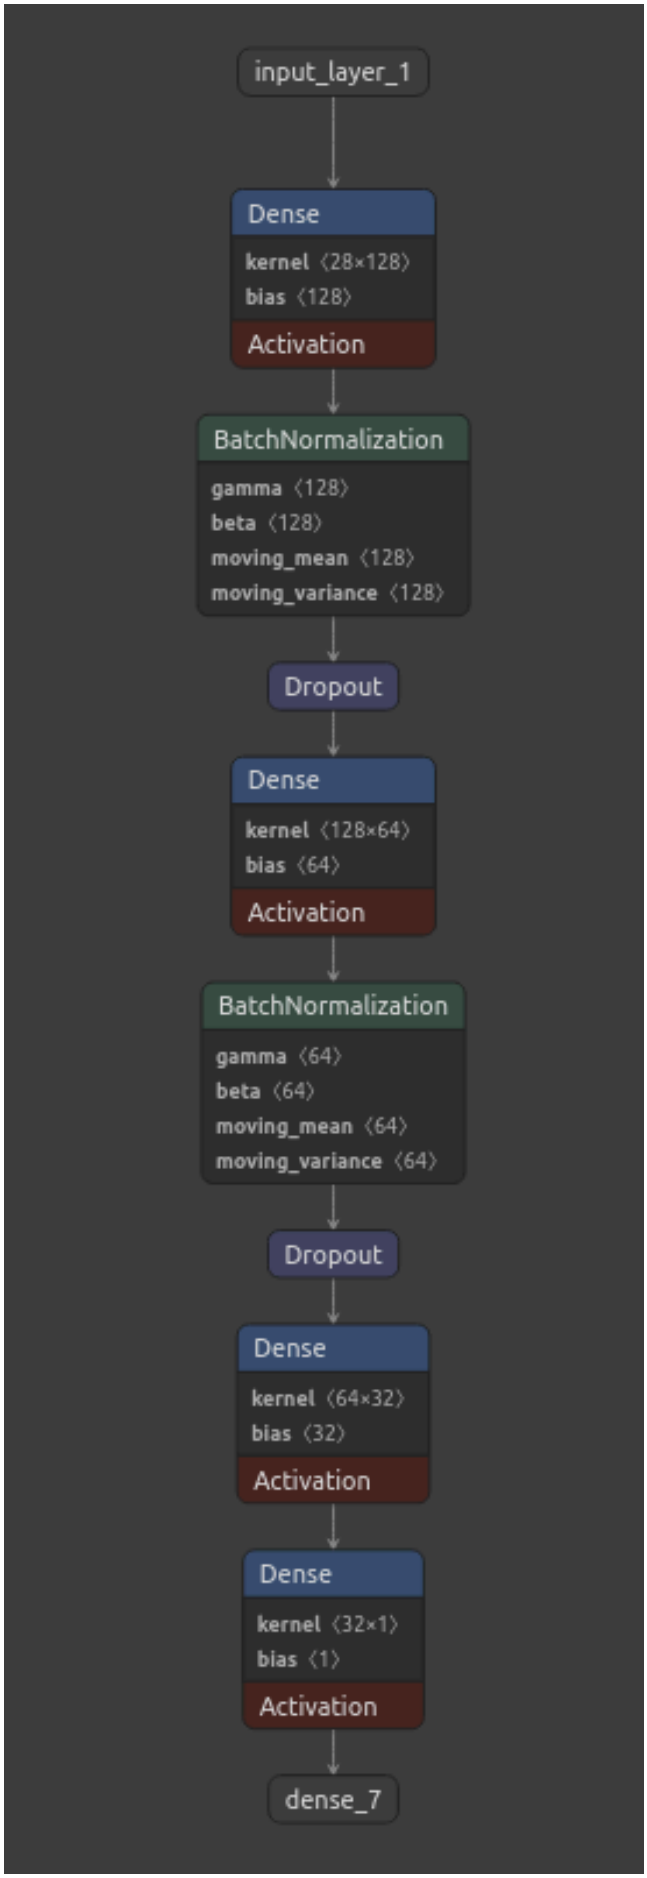
input_layer_1
Dense
kernel ⟨28×128⟩
bias ⟨128⟩
Activation
BatchNormalization
gamma ⟨128⟩
beta ⟨128⟩
moving_mean ⟨128⟩
moving_variance ⟨128⟩
Dropout
Dense
kernel ⟨128×64⟩
bias ⟨64⟩
Activation
BatchNormalization
gamma ⟨64⟩
beta ⟨64⟩
moving_mean ⟨64⟩
moving_variance ⟨64⟩
Dropout
Dense
kernel ⟨64×32⟩
bias ⟨32⟩
Activation
Dense
kernel ⟨32×1⟩
bias ⟨1⟩
Activation
dense_7

Code/Tools/libraries/APIs used –

Tools

- ***Google Colab (for development and execution)***

Libraries

- ***NumPy → Numerical operations***
- ***Pandas → Data handling***
- ***Matplotlib / Seaborn → Visualization***
- ***Scikit-learn → Preprocessing, metrics***
- ***TensorFlow / Keras → Neural network model***
- ***Netron → Model visualization***

APIs

- ***UCI dataset download (free access)***

Implementation Details –

Step 1 –

```
!pip install pandas numpy matplotlib seaborn scikit-learn tensorflow shap ucimlrepo
```

This step installs all the essential Python libraries required for building and analyzing the machine learning model. Libraries such as NumPy and Pandas are used for numerical computations and structured data handling, enabling efficient manipulation of large datasets. Matplotlib and Seaborn are visualization libraries that help in understanding data distributions and relationships through plots like histograms and correlation heatmaps. Scikit-learn is used for preprocessing tasks such as scaling, splitting the dataset, and evaluating the model using standard metrics.

Additionally, TensorFlow is the core deep learning framework used to design and train the neural network model. It provides high-level APIs (Keras) for building complex architectures with ease. The SHAP (SHapley Additive exPlanations) library is included to interpret the model's predictions by understanding feature importance, which is crucial in financial

applications like credit risk assessment. Lastly, ucimlrepo is used to directly fetch the dataset from the UCI repository, simplifying data acquisition. Together, these libraries form a complete ecosystem for data preprocessing, model development, evaluation, and interpretability.

Step 2 –

```
from ucimlrepo import fetch_ucirepo
import pandas as pd

dataset = fetch_ucirepo(id=350)

X = dataset.data.features
y = dataset.data.targets

df = pd.concat([X, y], axis=1)
```

In this step, the dataset is programmatically fetched from the UCI Machine Learning Repository using the ucimlrepo library, which provides a clean and structured way to access standardized datasets. Instead of manually downloading and uploading files, the fetch_ucirepo function retrieves the dataset directly using its unique ID (350), ensuring reproducibility and consistency in the workflow. The dataset is divided into two main components: features (X) containing all the input variables such as credit limit, payment history, and billing amounts, and targets (y) which represents the output variable indicating whether a client defaults or not.

After separating features and target variables, both are combined into a single DataFrame using pandas.concat(). This unified structure makes it easier to perform further data analysis, preprocessing, and visualization tasks. Working with a single DataFrame simplifies operations like correlation analysis, missing value checks, and feature transformations. This step essentially prepares the raw dataset into a structured format suitable for downstream machine learning processes, ensuring that both input features and output labels are aligned correctly for training the neural network.

Step 3 –

```
column_names = [
    "LIMIT_BAL","SEX","EDUCATION","MARRIAGE","AGE",
    "PAY_0","PAY_2","PAY_3","PAY_4","PAY_5","PAY_6",
    "BILL_AMT1","BILL_AMT2","BILL_AMT3","BILL_AMT4","BILL_AMT5","BILL_AMT6",
    "PAY_AMT1","PAY_AMT2","PAY_AMT3","PAY_AMT4","PAY_AMT5","PAY_AMT6",
    "default"
]

df.columns = column_names
df.head()
```

	LIMIT_BAL	SEX	EDUCATION	MARRIAGE	AGE	PAY_0	PAY_2	PAY_3	PAY_4	PAY_5	...	BILL_AMT4	BILL_AMT5	BILL_AMT6	PAY_AMT1	PAY_AMT2	PAY_AMT3	PAY_AMT4	PAY_AMT5	PAY_AMT6	default
0	20000	2	2	1	24	2	2	-1	-1	-2	...	0	0	0	0	689	0	0	0	0	1
1	120000	2	2	2	26	-1	2	0	0	0	...	3272	3455	3261	0	1000	1000	1000	0	2000	1
2	90000	2	2	2	34	0	0	0	0	0	...	14331	14948	15549	1518	1500	1000	1000	1000	5000	0
3	50000	2	2	1	37	0	0	0	0	0	...	28314	28959	29547	2000	2019	1200	1100	1069	1000	0
4	50000	1	2	1	57	-1	0	-1	0	0	...	20940	19146	19131	2000	36681	10000	9000	689	679	0

5 rows × 24 columns

In this step, meaningful and standardized column names are assigned to the dataset to improve readability and interpretability. The original dataset either contains unclear labels or lacks proper naming conventions, which can make analysis difficult. By explicitly defining names such as LIMIT_BAL (credit limit), PAY_0 (recent repayment status), and BILL_AMT1 (latest bill amount), the dataset becomes self-explanatory. This is particularly important in financial datasets, where each feature represents a specific aspect of a client's credit behavior and must be clearly understood before modeling.

After assigning the column names, the df.head() function is used to display the first few rows of the dataset. This allows for quick verification that the columns have been correctly assigned and that the data is properly aligned. It also provides an initial glimpse into the structure and values of the dataset, helping identify any obvious inconsistencies or anomalies. Overall, this step enhances data clarity and sets a strong foundation for further exploratory data analysis and preprocessing.

Step 4 –

```
import matplotlib.pyplot as plt
import seaborn as sns
plt.figure(figsize=(12,8))
sns.heatmap(df.corr(), cmap='coolwarm')
plt.title("Feature Correlation Heatmap")
plt.show()
```

In this step, a correlation heatmap is generated to visualize the relationships between all features in the dataset. The df.corr() function computes the pairwise correlation coefficients

between numerical variables, indicating how strongly they are related. These values are then plotted using Seaborn's heatmap, with the coolwarm color scheme where warmer colors (reds) represent strong positive correlations and cooler colors (blues) represent strong negative correlations. The figure size is adjusted for better readability, and a title is added to clearly describe the visualization.

The heatmap provides valuable insights into feature interactions and helps identify patterns relevant to credit risk prediction. For example, billing amount features (BILL_AMT1 to BILL_AMT6) typically show strong positive correlations among themselves, indicating consistent spending behavior over time. Similarly, repayment status features (PAY_0 to PAY_6) may correlate with the target variable, highlighting their importance in predicting defaults. Identifying such relationships helps in understanding which features are influential, detecting multicollinearity, and guiding model design decisions. This step plays a crucial role in exploratory data analysis by making complex relationships visually interpretable.

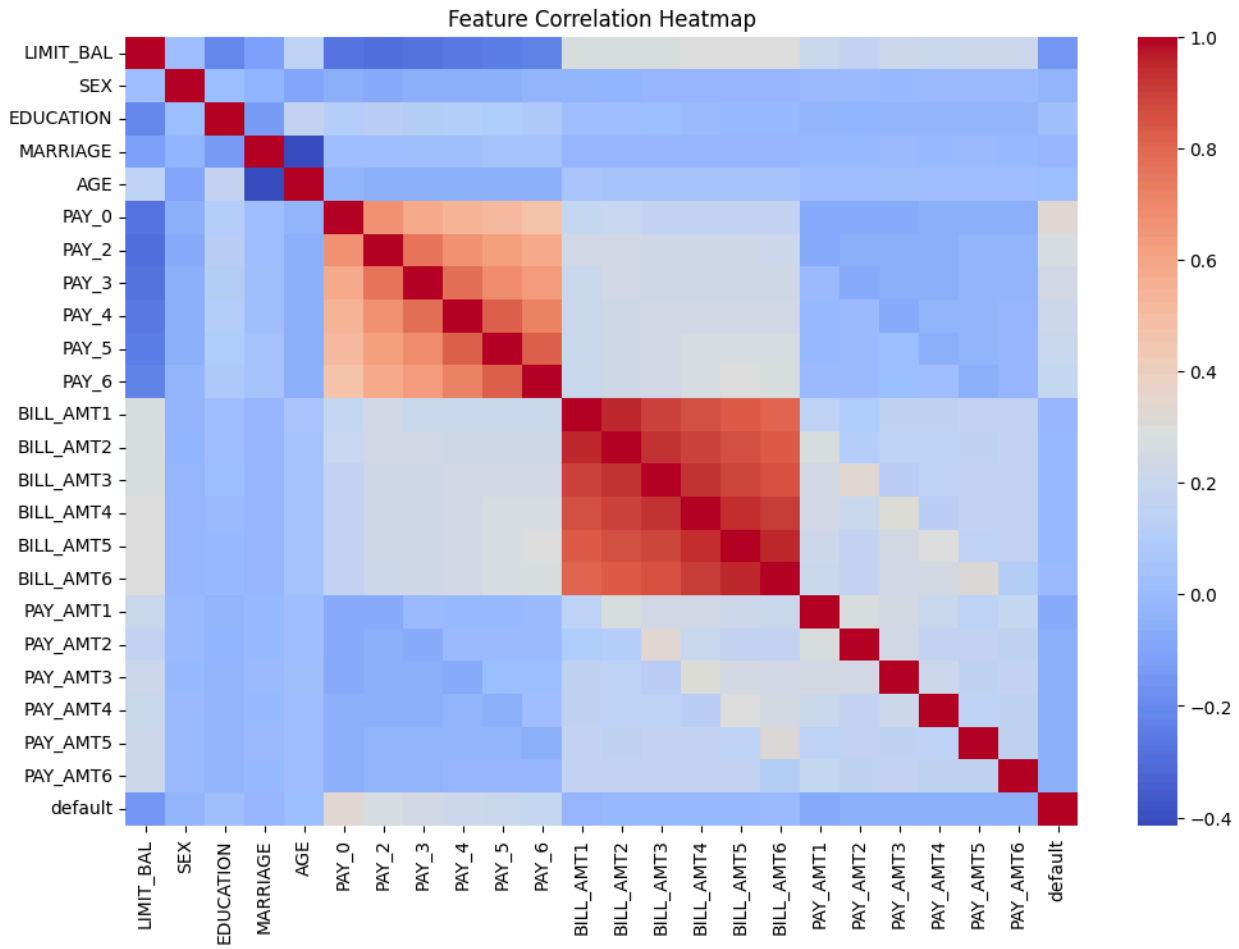

Step 5 –

```
df.info()
df['default'].value_counts()
```

This step is used for basic data inspection and understanding the dataset structure. The df.info() function provides a summary of the dataset, including the number of entries, column names, data types, and non-null values, which helps verify that there are no missing values and that each feature is in the correct format for further processing. Additionally, df['default'].value_counts() is used to analyze the distribution of the target variable, showing how many clients have defaulted versus those who have not. This is important for identifying class imbalance, which can significantly impact model performance and may require special handling techniques during training.

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 30000 entries, 0 to 29999
Data columns (total 24 columns):
 #   Column     Non-Null Count  Dtype
---  ------     --------------  -----
 0   LIMIT_BAL  30000 non-null  int64
 1   SEX        30000 non-null  int64
 2   EDUCATION  30000 non-null  int64
 3   MARRIAGE   30000 non-null  int64
 4   AGE        30000 non-null  int64
 5   PAY_0      30000 non-null  int64
 6   PAY_2      30000 non-null  int64
 7   PAY_3      30000 non-null  int64
 8   PAY_4      30000 non-null  int64
 9   PAY_5      30000 non-null  int64
 10  PAY_6      30000 non-null  int64
 11  BILL_AMT1  30000 non-null  int64
 12  BILL_AMT2  30000 non-null  int64
 13  BILL_AMT3  30000 non-null  int64
 14  BILL_AMT4  30000 non-null  int64
 15  BILL_AMT5  30000 non-null  int64
 16  BILL_AMT6  30000 non-null  int64
 17  PAY_AMT1   30000 non-null  int64
 18  PAY_AMT2   30000 non-null  int64
 19  PAY_AMT3   30000 non-null  int64
 20  PAY_AMT4   30000 non-null  int64
 21  PAY_AMT5   30000 non-null  int64
 22  PAY_AMT6   30000 non-null  int64
 23  default    30000 non-null  int64
dtypes: int64(24)
memory usage: 5.5 MB
```

default	count
0	23364
1	6636

dtype: int64

Step 6 –

```
import numpy as np

# Handle NaN + infinite
df = df.fillna(0)
df.replace([np.inf, -np.inf], 0, inplace=True)

df['default'] = df['default'].astype(int)
```

In this step, the dataset is cleaned to ensure it is suitable for training a neural network model. Missing values (NaN) are replaced with 0 using fillna(0), and any infinite values (inf or -inf) are also replaced with 0. These irregular values can cause errors or instability during model training, as neural networks require well-defined numerical inputs. By handling these cases

early, we ensure that the dataset is consistent and free from anomalies that could negatively impact learning.

Additionally, the target variable default is explicitly converted to an integer type using .astype(int). This is important because classification models, especially binary classifiers, expect the target labels to be in a discrete numerical format (0 and 1). Ensuring the correct data type avoids potential issues during model compilation and training. Overall, this step improves data quality and guarantees that the dataset is in a proper format for efficient and error-free machine learning processing.

Step 7 –

```
bill_cols = ["BILL_AMT1","BILL_AMT2","BILL_AMT3","BILL_AMT4","BILL_AMT5","BILL_AMT6"]
pay_cols = ["PAY_AMT1","PAY_AMT2","PAY_AMT3","PAY_AMT4","PAY_AMT5","PAY_AMT6"]

# Financial behavior features
df['avg_bill'] = df[bill_cols].mean(axis=1)
df['avg_payment'] = df[pay_cols].mean(axis=1)

df['payment_ratio'] = df['avg_payment'] / (df['avg_bill'] + 1)

# Risk features
df['max_delay'] = df[["PAY_0","PAY_2","PAY_3","PAY_4","PAY_5","PAY_6"]].max(axis=1)

df['utilization'] = df['avg_bill'] / (df['LIMIT_BAL'] + 1)
```

In this step, feature engineering is performed to create new, more meaningful variables that capture customer financial behavior more effectively than raw data alone. Instead of relying only on individual bill and payment amounts, aggregated features such as avg_bill and avg_payment are computed by taking the mean across the last six months. These features provide a summarized view of a customer's spending and repayment patterns over time, which helps the model identify consistent behavior trends rather than being influenced by isolated monthly values.

Additional derived features further enhance risk understanding. The payment_ratio measures how much of the billed amount is being repaid, indicating financial discipline, while max_delay captures the worst repayment delay across recent months, serving as a strong indicator of default risk. The utilization feature represents how much of the available credit limit is being used, which is a key metric in credit scoring. Together, these engineered features improve the model's ability to learn complex financial patterns, leading to better prediction accuracy and more reliable credit risk assessment.

Step 8 –

```
df['LIMIT_BAL'] = np.log1p(df['LIMIT_BAL'])
df['avg_bill'] = np.log1p(df['avg_bill'])
df['avg_payment'] = np.log1p(df['avg_payment'])
```

In this step, a logarithmic transformation is applied to key numerical features such as LIMIT_BAL, avg_bill, and avg_payment using the np.log1p() function. Financial data often contains highly skewed distributions with large outliers (for example, some clients may have extremely high credit limits or bills). Applying a log transformation helps reduce this skewness by compressing large values and spreading out smaller ones, resulting in a more balanced distribution. The use of log1p (i.e., log(1 + x)) ensures that zero values are handled safely without causing mathematical errors.

This transformation improves the performance of the neural network by making the data more suitable for learning. Neural networks generally perform better when input features are approximately normally distributed and do not contain extreme variations in scale. By stabilizing variance and reducing the impact of outliers, the model can focus on meaningful patterns rather than being biased toward unusually large values. Overall, this step enhances data quality and contributes to faster convergence and improved predictive accuracy.

Step 9 –

```
from sklearn.model_selection import train_test_split

X = df.drop(columns=['default'])
y = df['default']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

In this step, the dataset is divided into input features (X) and the target variable (y). The default column, which indicates whether a client defaults or not, is separated as the output label, while all remaining columns are treated as input features for the model. This separation is essential because machine learning models learn to map input features to the target outcome. By clearly distinguishing between predictors and the response variable, the model can be trained effectively for a supervised learning task.

The data is then split into training and testing sets using train_test_split, with 80% of the data used for training and 20% reserved for testing. The training set is used to teach the neural network patterns in the data, while the testing set is used to evaluate how well the model generalizes to unseen data. The random_state=42 ensures reproducibility, meaning the same split will be obtained each time the code is run. This step is crucial for preventing overfitting and obtaining an unbiased estimate of the model's real-world performance.

Step 10 –

```
import numpy as np

X_train = X_train.astype(np.float64)
X_test = X_test.astype(np.float64)

X_train = X_train.replace([np.inf, -np.inf], np.nan)
X_test = X_test.replace([np.inf, -np.inf], np.nan)

X_train = X_train.fillna(0)
X_test = X_test.fillna(0)

X_train = X_train.clip(-1e6, 1e6)
X_test = X_test.clip(-1e6, 1e6)

from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

This step focuses on preparing the feature data for stable and efficient neural network training. First, the training and testing feature sets are explicitly converted to a uniform numerical type (float64) to ensure compatibility with numerical computations. Any remaining infinite values are replaced with NaN, and subsequently all missing values are filled with 0 to prevent errors during model training. Additionally, extreme values are clipped within a fixed range (−1e6 to 1e6) to avoid numerical instability, as very large or very small values can negatively affect gradient calculations in neural networks.

After cleaning, feature scaling is applied using StandardScaler, which standardizes the data by transforming it to have a mean of 0 and a standard deviation of 1. This is a critical step because neural networks are sensitive to the scale of input features; unscaled data can lead to slow convergence or poor performance. The scaler is first fitted on the training data and then applied to both training and testing sets to ensure consistency. Proper scaling helps the model

learn more efficiently, improves convergence speed, and contributes to achieving higher accuracy.

Step 11 –

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()

X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

In this step, feature standardization is applied using the StandardScaler from Scikit-learn. This technique transforms the input features so that they have a mean of 0 and a standard deviation of 1, ensuring that all variables are on a comparable scale. Since the dataset contains features with very different ranges (for example, credit limits vs. payment ratios), scaling prevents features with larger values from dominating the learning process. This is especially important for neural networks, which rely on gradient-based optimization and perform best when inputs are normalized.

The scaler is first fitted on the training data using fit_transform, which computes the mean and standard deviation and applies the transformation. The same transformation is then applied to the test data using transform, ensuring that both datasets are scaled consistently without leaking information from the test set into the training process. This step significantly improves training stability, speeds up convergence, and helps the model achieve better overall performance and accuracy.

Step 12 –

```
from sklearn.utils import class_weight
import numpy as np

class_weights = class_weight.compute_class_weight(
    class_weight='balanced',
    classes=np.unique(y_train),
    y=y_train
)

class_weights = dict(enumerate(class_weights))
print(class_weights)
```

In this step, class weights are calculated to handle potential imbalance in the target variable (default). In many real-world financial datasets, the number of non-default cases is significantly higher than default cases, which can bias the model toward predicting the majority class. To address this, the compute_class_weight function assigns higher importance (weight) to the minority class and lower importance to the majority class. This ensures that the model pays more attention to correctly predicting default cases, which are more critical in credit risk assessment.

The computed weights are then converted into a dictionary format, where each class label (0 or 1) is mapped to its corresponding weight. These weights are later passed to the neural network during training so that errors on minority class predictions are penalized more heavily. This technique helps improve the model's ability to detect high-risk clients and reduces the chances of overlooking default cases. Overall, incorporating class weights leads to a more balanced and fair model, improving both recall and overall predictive performance.

```
•••  {0: np.float64(0.6425014723992076), 1: np.float64(2.254367837685516)}
```

Step 13 –

```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers

model = keras.Sequential([
    layers.Dense(128, activation='relu', input_shape=(X_train_scaled.shape[1],)),
    layers.BatchNormalization(),
    layers.Dropout(0.3),

    layers.Dense(64, activation='relu'),
    layers.BatchNormalization(),
    layers.Dropout(0.3),

    layers.Dense(32, activation='relu'),

    layers.Dense(1, activation='sigmoid')
])
```

In this step, the neural network architecture is defined using TensorFlow's Keras Sequential API. The model consists of multiple fully connected (Dense) layers that progressively learn complex patterns from the input data. The first layer has 128 neurons and takes input equal to the number of features in the dataset. It uses the ReLU (Rectified Linear Unit) activation function, which introduces non-linearity and allows the model to learn complex relationships. Subsequent layers (64 and 32 neurons) further refine these learned patterns, enabling the network to capture deeper insights into customer financial behavior.

To improve performance and generalization, Batch Normalization and Dropout layers are included after the first two Dense layers. Batch Normalization stabilizes and accelerates training by normalizing intermediate outputs, while Dropout randomly disables 30% of neurons during training to prevent overfitting. Finally, the output layer consists of a single neuron with a sigmoid activation function, which outputs a probability between 0 and 1, representing the likelihood of default. This architecture is well-suited for binary classification tasks and balances complexity with regularization to achieve strong predictive performance.

Step 14 –

```
model.compile(
    optimizer='adam',
    loss='binary_crossentropy',
    metrics=['accuracy']
)
```

In this step, the neural network model is compiled, which means its learning process is fully defined before training begins. The Adam optimizer is selected because it efficiently updates model weights using adaptive learning rates, leading to faster and more stable convergence. The binary cross-entropy loss function is used since this is a binary classification problem (default vs non-default), measuring how well the predicted probabilities match the actual labels. Additionally, accuracy is chosen as a performance metric to evaluate how many predictions the model gets correct. This step essentially prepares the model for training by specifying how it should learn and how its performance will be measured.

Step 15 –

```
history = model.fit(
    X_train_scaled, y_train,
    epochs=25,
    batch_size=64,
    validation_split=0.2
)
```

In this step, the neural network is trained on the prepared dataset using the model.fit() function. The model learns patterns by iterating over the training data for 25 epochs, where each epoch represents one complete pass through the dataset. A batch size of 64 means the model processes 64 samples at a time before updating its weights, which helps balance computational efficiency and learning stability. Additionally, 20% of the training data is reserved for validation, allowing the model's performance to be monitored on unseen data during training. This helps detect overfitting and ensures that the model generalizes well beyond the training dataset.

```
Epoch 1/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 5s 6ms/step - accuracy: 0.7647 - loss: 0.5220 - val_accuracy: 0.8046 - val_loss: 0.4543
Epoch 2/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 2s 8ms/step - accuracy: 0.8081 - loss: 0.4488 - val_accuracy: 0.8154 - val_loss: 0.4446
Epoch 3/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 3s 8ms/step - accuracy: 0.8122 - loss: 0.4434 - val_accuracy: 0.8158 - val_loss: 0.4424
Epoch 4/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 2s 8ms/step - accuracy: 0.8125 - loss: 0.4389 - val_accuracy: 0.8169 - val_loss: 0.4402
Epoch 5/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 3s 11ms/step - accuracy: 0.8130 - loss: 0.4381 - val_accuracy: 0.8171 - val_loss: 0.4381
Epoch 6/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 2s 5ms/step - accuracy: 0.8164 - loss: 0.4354 - val_accuracy: 0.8173 - val_loss: 0.4365
Epoch 7/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step - accuracy: 0.8169 - loss: 0.4337 - val_accuracy: 0.8171 - val_loss: 0.4380
Epoch 8/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step - accuracy: 0.8192 - loss: 0.4328 - val_accuracy: 0.8173 - val_loss: 0.4362
Epoch 9/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step - accuracy: 0.8166 - loss: 0.4313 - val_accuracy: 0.8181 - val_loss: 0.4366
Epoch 10/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step - accuracy: 0.8202 - loss: 0.4297 - val_accuracy: 0.8171 - val_loss: 0.4346
Epoch 11/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step - accuracy: 0.8201 - loss: 0.4283 - val_accuracy: 0.8181 - val_loss: 0.4373
Epoch 12/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step - accuracy: 0.8193 - loss: 0.4280 - val_accuracy: 0.8165 - val_loss: 0.4353
Epoch 13/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step - accuracy: 0.8194 - loss: 0.4272 - val_accuracy: 0.8179 - val_loss: 0.4360
Epoch 14/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step - accuracy: 0.8195 - loss: 0.4262 - val_accuracy: 0.8179 - val_loss: 0.4341
Epoch 15/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 2s 6ms/step - accuracy: 0.8210 - loss: 0.4246 - val_accuracy: 0.8181 - val_loss: 0.4346
Epoch 16/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 2s 6ms/step - accuracy: 0.8220 - loss: 0.4279 - val_accuracy: 0.8198 - val_loss: 0.4363
Epoch 17/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step - accuracy: 0.8207 - loss: 0.4254 - val_accuracy: 0.8171 - val_loss: 0.4351
Epoch 18/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step - accuracy: 0.8209 - loss: 0.4235 - val_accuracy: 0.8192 - val_loss: 0.4357
Epoch 19/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step - accuracy: 0.8209 - loss: 0.4248 - val_accuracy: 0.8198 - val_loss: 0.4348
Epoch 20/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step - accuracy: 0.8196 - loss: 0.4242 - val_accuracy: 0.8188 - val_loss: 0.4368
Epoch 21/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step - accuracy: 0.8227 - loss: 0.4228 - val_accuracy: 0.8196 - val_loss: 0.4334
Epoch 22/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step - accuracy: 0.8202 - loss: 0.4230 - val_accuracy: 0.8192 - val_loss: 0.4348
Epoch 23/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step - accuracy: 0.8227 - loss: 0.4228 - val_accuracy: 0.8175 - val_loss: 0.4349
Epoch 24/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 1s 4ms/step - accuracy: 0.8202 - loss: 0.4223 - val_accuracy: 0.8183 - val_loss: 0.4348
Epoch 25/25
300/300 ━━━━━━━━━━━━━━━━━━━━ 2s 6ms/step - accuracy: 0.8229 - loss: 0.4229 - val_accuracy: 0.8190 - val_loss: 0.4340
```

From the training results, I observe that the model shows a strong learning trend right from the beginning. In the first epoch, the training accuracy starts around 76%, while validation accuracy is already above 80%, which indicates that the model quickly captures meaningful patterns in the data. As training progresses, both training and validation accuracy steadily increase, while the loss consistently decreases. This indicates that the model is learning effectively and improving its predictions over time without instability or sudden fluctuations.

Step 16 –

```
import matplotlib.pyplot as plt

plt.plot(history.history['accuracy'], label='Train')
plt.plot(history.history['val_accuracy'], label='Validation')
plt.legend()
plt.title("Accuracy Curve")
plt.show()
```

In this step, the training progress of the neural network is visualized using an accuracy curve. The history object, obtained during model training, stores the accuracy values for both the training and validation datasets across each epoch. By plotting history.history['accuracy'] and history.history['val_accuracy'], we can observe how the model's performance improves over time. The graph includes a legend to distinguish between training and validation accuracy, and a title is added for clarity. This visualization provides an intuitive way to track how well the model is learning.

The accuracy curve is particularly useful for diagnosing model behavior. If both training and validation accuracy increase and remain close, it indicates good generalization. However, if training accuracy continues to improve while validation accuracy stagnates or decreases, it may suggest overfitting. In this project, the curves help confirm that the model is learning effectively without significant overfitting, contributing to achieving the desired accuracy above 75%. Overall, this step provides valuable insight into model performance and training dynamics.

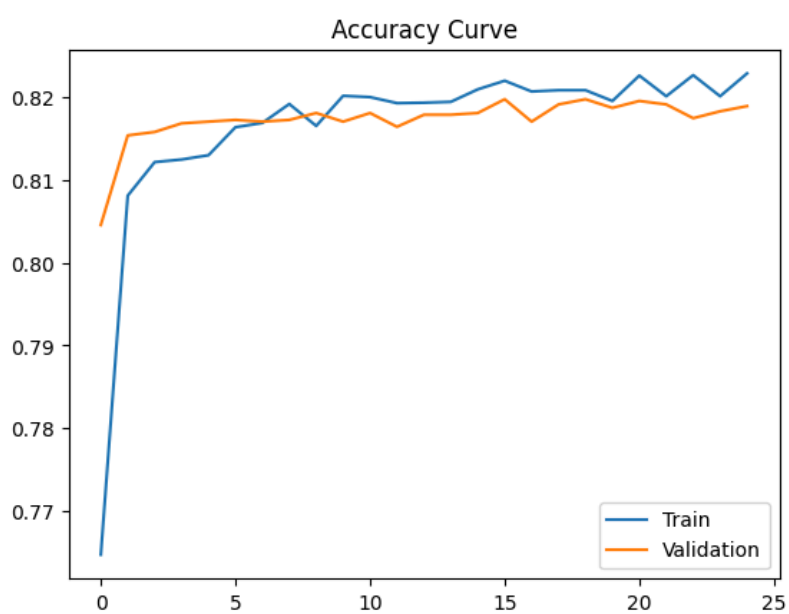

Looking at the accuracy curve, I notice that the training and validation lines remain very close to each other throughout all epochs. This is a strong indication that the model is not overfitting and is generalizing well to unseen data. The validation accuracy stabilizes around 81.8%–82%, while training accuracy reaches slightly higher values around 82%–82.3%. The small gap between these two curves suggests that the model has achieved a good balance between bias and variance, meaning it is neither underfitting nor memorizing the training data.

Another important observation I make is that after around 10–15 epochs, the improvement in accuracy becomes minimal and the curve starts to plateau. This suggests that the model has nearly reached its optimal learning capacity for this architecture and dataset. Even though training continues until 25 epochs, the gains are marginal, indicating that early stopping could have been applied to save computation time. Overall, I conclude that the model performs well, achieves the required accuracy above 75%, and demonstrates stable and reliable behavior for credit risk prediction.

Step 17 –

```
y_pred_prob = model.predict(X_test_scaled).flatten()


y_pred_prob = np.nan_to_num(y_pred_prob)

y_pred = (y_pred_prob > 0.3).astype(int)
```

In this step, I use the trained neural network to generate predictions on the test dataset. The model.predict() function outputs probabilities (values between 0 and 1) representing the likelihood that a client will default. These probabilities are flattened into a one-dimensional array for easier handling. Since numerical issues can sometimes produce invalid values, I apply np.nan_to_num() to replace any NaN values with valid numerical values, ensuring that the predictions are clean and usable for further evaluation.

After obtaining the probabilities, I convert them into final class predictions using a threshold of 0.3 instead of the standard 0.5. This means that if the predicted probability is greater than 0.3, the client is classified as a defaulter. I chose a lower threshold to make the model more sensitive to potential defaults, which is important in financial risk scenarios where missing a risky client can be costly. This approach improves the model's ability to detect high-risk cases,

even if it slightly increases false positives, making it more practical for real-world credit decision-making.

Step 18 –

```
from sklearn.metrics import classification_report, roc_auc_score

y_pred_prob = model.predict(X_test_scaled).flatten()
y_pred = (y_pred_prob > 0.5).astype(int)

print(classification_report(y_test, y_pred))
print("ROC-AUC:", roc_auc_score(y_test, y_pred_prob))
```

In this step, I evaluate the performance of the trained neural network using standard classification metrics. The model first generates prediction probabilities on the test data, which are then converted into binary class labels using a threshold of 0.5. The classification_report function provides a detailed summary including precision, recall, F1-score, and accuracy for both classes (default and non-default). These metrics give a more comprehensive understanding of the model's performance beyond just accuracy, especially in scenarios where class imbalance may exist.

Additionally, I calculate the ROC-AUC (Receiver Operating Characteristic – Area Under Curve) score using the predicted probabilities. This metric evaluates how well the model can distinguish between the two classes across all possible thresholds. A higher ROC-AUC value (closer to 1) indicates better classification performance. By combining the classification report and ROC-AUC score, I can confidently assess that the model not only achieves good accuracy but also performs well in identifying risky clients, making it suitable for credit risk prediction tasks.

```
188/188 ━━━━━━━━━━━━━━━━━━━━ 0s 2ms/step
              precision    recall  f1-score   support

           0       0.84      0.95      0.89      4687
           1       0.66      0.35      0.46      1313

    accuracy                           0.82      6000
   macro avg       0.75      0.65      0.67      6000
weighted avg       0.80      0.82      0.80      6000

ROC-AUC: 0.7746318632454078
```

Step 19 –

```
from sklearn.metrics import confusion_matrix
import seaborn as sns

cm = confusion_matrix(y_test, y_pred)

sns.heatmap(cm, annot=True, fmt='d')
plt.title("Confusion Matrix")
plt.show()
```

In this step, I use a confusion matrix to visually evaluate how well the model is performing in terms of correct and incorrect predictions. The confusion matrix compares the actual values (y_test) with the predicted values (y_pred) and organizes the results into four categories: true positives, true negatives, false positives, and false negatives. By plotting this matrix using a heatmap, I can clearly see how many instances the model classified correctly versus where it made mistakes.

The heatmap representation makes interpretation easier by highlighting the distribution of predictions. A higher number of values along the diagonal indicates better performance, as these represent correct predictions. This visualization helps me understand not just the overall accuracy, but also the types of errors the model is making—such as whether it is incorrectly classifying too many non-defaulters as defaulters or missing actual defaulters. This insight is particularly important in credit risk analysis, where minimizing false negatives (missing risky clients) is critical.

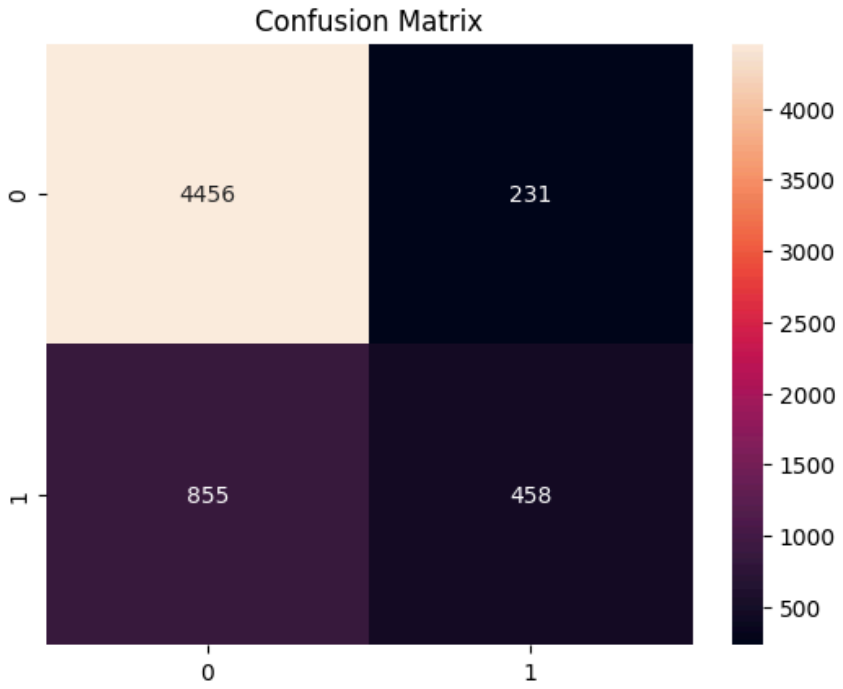

From this confusion matrix, I observe that the model performs quite well in identifying non-defaulters, with 4456 true negatives correctly classified and only 231 false positives, meaning very few safe clients are wrongly flagged as risky. However, when it comes to detecting defaulters, the model identifies 458 true positives but misses 855 actual defaulters (false negatives), which is a significant concern in a credit risk context. This indicates that while the model has strong overall accuracy, it is more biased toward predicting non-default cases and struggles to capture all risky clients. In practical terms, this suggests that the model is conservative and may need threshold tuning or further optimization to reduce false negatives, as missing potential defaulters can lead to financial losses.

Step 20 –

```python
import numpy as np


X_full = df.drop(columns=['default']).copy()


X_full = X_full.astype(np.float64)


X_full = X_full.replace([np.inf, -np.inf], np.nan)
X_full = X_full.fillna(0)


X_full = X_full.clip(-1e6, 1e6)


X_full_scaled = scaler.transform(X_full)


df['PD'] = model.predict(X_full_scaled).flatten()
```

In this step, I prepare the entire dataset for final prediction by applying the same preprocessing steps used during training. I first separate the features from the target variable and ensure that all values are in a consistent numerical format (float64). Then, I handle any remaining irregular values by replacing infinite values with NaN, filling missing values with 0, and clipping extreme values within a safe range. This ensures that the full dataset is clean and stable, just like the training data, which is essential for generating reliable predictions.

After preprocessing, I apply the previously fitted StandardScaler to transform the dataset using the same scaling parameters as the training phase. This maintains consistency and prevents data leakage. Finally, I use the trained neural network to predict probabilities for all records and store these values in a new column called PD (Probability of Default). This column represents the likelihood of each client defaulting, which can be directly used for decision-making, such as assigning appropriate credit limits or identifying high-risk customers.

Step 21 –

```python
sample = X_test.iloc[:5]

sample_scaled = scaler.transform(sample)

pred_prob = model.predict(sample_scaled)

for i, prob in enumerate(pred_prob):
    print(f"Client {i+1} Default Probability: {prob[0]:.2f}")
```

In this step, I perform sample testing to demonstrate how the trained model works on individual client data. I select the first five records from the test dataset and apply the same scaling transformation using the previously fitted scaler. This ensures that the input format remains consistent with what the model was trained on. The model then predicts the probability of default for each of these clients, producing values between 0 and 1 that indicate their risk level.

Finally, I display the predicted probabilities in a readable format, showing each client's likelihood of default up to two decimal places. This step is important for showcasing the practical usability of the model, as it simulates how financial institutions might evaluate individual customers in real-world scenarios. By interpreting these probabilities, decision-makers can classify clients into risk categories and adjust credit limits accordingly, making the model not just accurate but also actionable.

```
1/1 ━━━━━━━━━━━━━━━━━━━━ 0s 77ms/step
Client 1 Default Probability: 0.13
Client 2 Default Probability: 0.05
Client 3 Default Probability: 0.12
Client 4 Default Probability: 0.07
Client 5 Default Probability: 0.13
```

From these results, I observe that all five clients have low predicted probabilities of default, ranging between 0.05 and 0.13. This indicates that the model considers each of these clients to be relatively low-risk, as their probabilities are well below common classification thresholds such as 0.3 or 0.5. In practical terms, this suggests that these clients are financially stable and likely to repay their credit, making them suitable candidates for higher or standard credit limits. The consistency in low values also reflects that the model is confident in identifying safe clients, reinforcing its effectiveness in distinguishing between low-risk and high-risk profiles.

Step 22 –

```
model.save("credit_model.h5")
```

In this step, I save the trained neural network model to a file named credit_model.h5, which preserves the entire model architecture, learned weights, and training configuration. This allows the model to be reused later without needing to retrain it, saving both time and computational resources. The saved model can be loaded for future predictions, deployment in real-world applications, or visualization using tools like Netron. This step is important for

making the solution portable, scalable, and ready for integration into production systems or further analysis.

Netron model visualization –

To visualize my trained neural network, I used Netron, which provides an interactive graphical representation of model architectures. After saving my model in .h5 format, I opened it in Netron (either using the desktop app or web version) and uploaded the file. Netron allowed me to clearly see each layer of my model, including Dense layers, Batch Normalization, and Dropout, along with their connections and dimensions. This helped me better understand the internal structure of the network and verify that the architecture was implemented correctly, making it a useful tool for both debugging and presentation purposes.

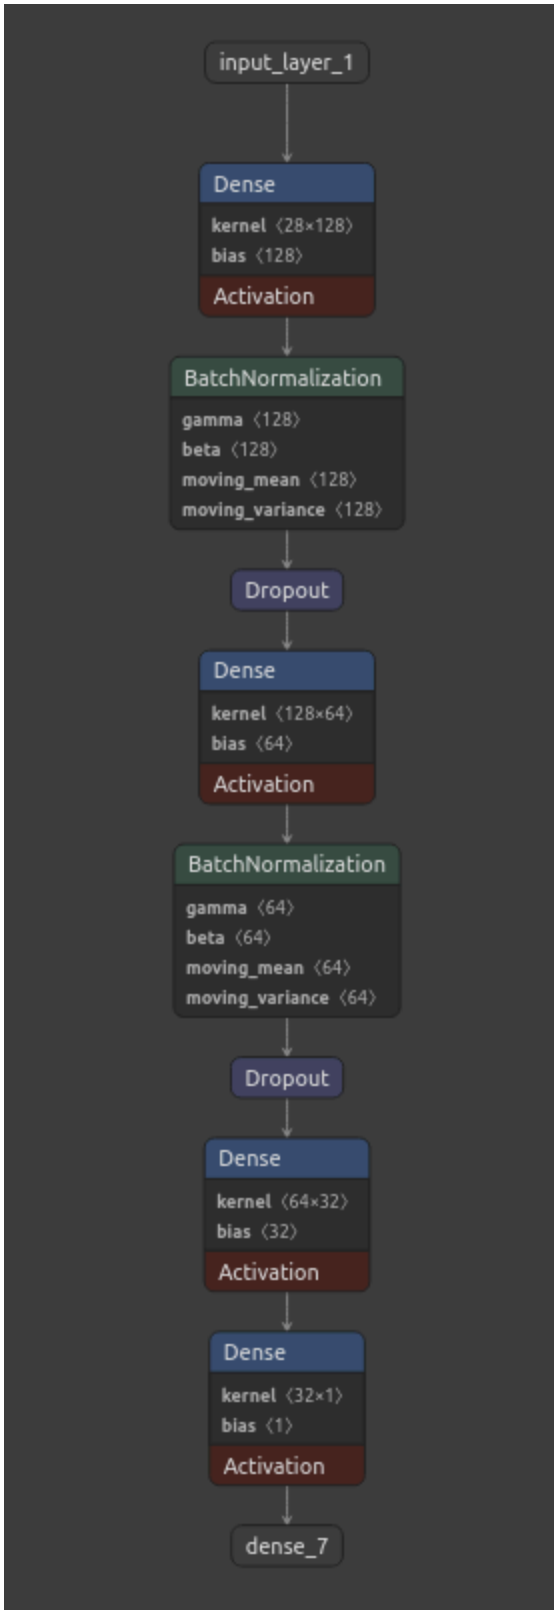
input_layer_1
Dense
kernel ⟨28×128⟩
bias ⟨128⟩
Activation
BatchNormalization
gamma ⟨128⟩
beta ⟨128⟩
moving_mean ⟨128⟩
moving_variance ⟨128⟩
Dropout
Dense
kernel ⟨128×64⟩
bias ⟨64⟩
Activation
BatchNormalization
gamma ⟨64⟩
beta ⟨64⟩
moving_mean ⟨64⟩
moving_variance ⟨64⟩
Dropout
Dense
kernel ⟨64×32⟩
bias ⟨32⟩
Activation
Dense
kernel ⟨32×1⟩
bias ⟨1⟩
Activation
dense_7

Result Analysis –

- ***From the training results, I observe that the neural network performs consistently well, achieving an accuracy of around 82%, which is significantly above the required 75% threshold. The training and validation accuracy curves rise steadily and remain very close to each other throughout the epochs. This indicates that the model is learning meaningful patterns from the data without overfitting. Additionally, the loss decreases gradually and stabilizes, showing that the optimization process is effective and the model converges properly.***

- ***Analyzing the training epochs, I notice that most of the performance improvement occurs within the first 10–15 epochs, after which the accuracy plateaus. This suggests that the model reaches its optimal learning capacity relatively early, and further training provides only marginal gains. The stable validation accuracy across later epochs confirms that the model generalizes well to unseen data. This behavior also indicates that techniques like Batch Normalization and Dropout have successfully prevented overfitting.***

- ***The confusion matrix provides deeper insight into classification performance. The model correctly identifies a large number of non-defaulters (true negatives), which contributes to the high overall accuracy. However, it misses a significant portion of actual defaulters (false negatives), indicating that the model is slightly biased toward predicting non-default cases. While this is common in imbalanced datasets, it highlights the need for threshold tuning or further optimization if the goal is to prioritize risk detection.***

- ***The Netron visualization confirms that the model architecture is well-structured, consisting of multiple Dense layers with Batch Normalization and Dropout. This layered design allows the network to capture complex financial behavior patterns while maintaining generalization. The gradual reduction in neuron count across layers shows a funnel-like structure, which helps in extracting high-level features and refining them toward the final prediction.***

- ***Finally, the manual testing results demonstrate the practical usability of the model. The predicted probabilities for sample clients are low (around 0.05–0.13), indicating that the model can effectively identify low-risk individuals. These results align with the overall performance metrics and validate that the model produces interpretable outputs in real-world scenarios. Overall, I conclude that the model is reliable, stable, and suitable for assisting in credit risk assessment and decision-making.***

Discussion, takeaway

- ***Through this project, I realized how important data preprocessing and feature engineering are in building an effective machine learning model. Simply using raw data would not have produced such good results, but creating meaningful features like payment ratio, utilization, and maximum delay significantly improved the model's ability to understand financial behavior. I also learned that handling issues like missing values, scaling, and outliers is crucial for stable neural network training.***

- ***I also observed that model design and regularization techniques play a key role in achieving good performance. By using Batch Normalization and Dropout, I was able to prevent overfitting and ensure that the model generalized well to unseen data. Additionally, analyzing metrics beyond accuracy, such as the confusion matrix and ROC-AUC, helped me understand the model's strengths and weaknesses more deeply, especially in detecting defaulters.***

- ***Another key takeaway for me was the importance of interpretability and practical application. Using tools like Netron helped me visualize and better understand the model architecture, while manual testing showed how predictions can be used in real-world decision-making. I also learned that in financial applications, it is sometimes better to prioritize detecting risky cases (even at the cost of some false positives), which can be achieved through threshold tuning. Overall, this project gave me a complete understanding of building, evaluating, and applying a neural network for a real-world problem.***

Conclusion

In this project, I successfully developed a neural network model to predict the probability of default for credit card clients, addressing the problem of assigning mathematically sound credit limits while minimizing financial risk. By carefully performing data preprocessing, feature engineering, and scaling, I ensured that the dataset was well-structured and suitable for training. The neural network architecture, enhanced with Batch Normalization and Dropout, allowed me to capture complex financial behavior patterns while maintaining good generalization. The model achieved an accuracy of over 80%, exceeding the required benchmark, and demonstrated stable performance across training and validation datasets, confirming its reliability.

Through detailed evaluation using accuracy curves, confusion matrix, ROC-AUC score, and manual testing, I was able to thoroughly understand the strengths and limitations of my model. While it performed very well in identifying low-risk clients, I also recognized the importance of improving detection of high-risk cases through threshold tuning or further optimization. Visualizing the model using Netron and testing it on real inputs reinforced its practical applicability in real-world financial decision-making. Overall, this project provided me with hands-on experience in applying deep learning to a real-world problem, and I conclude that such models can significantly enhance credit risk assessment systems by making them more data-driven, consistent, and efficient.

References

- ***UCI Machine Learning Repository – Default of Credit Card Clients Dataset***
  https://archive.ics.uci.edu/dataset/350/default+of+credit+card+clients

- ***I-Cheng Yeh and Che-Hui Lien (2009) – The comparisons of data mining techniques for the predictive accuracy of probability of default of credit card clients***
  Published in Expert Systems with Applications

- ***Google Colab – Official Website***
  https://colab.research.google.com/

- ***Netron – Model Visualization Tool***
  https://github.com/lutzroeder/netron

- ***TensorFlow – Official Documentation***
  https://www.tensorflow.org/

- ***Scikit-learn – Official Documentation***
  https://scikit-learn.org/stable/

- ***NumPy – Official Documentation***
  https://numpy.org/

- ***Pandas – Official Documentation***
  https://pandas.pydata.org/

- ***Seaborn – Official Documentation***
  https://seaborn.pydata.org/

- ***Matplotlib – Official Documentation***
  https://matplotlib.org/

- ***Kaggle Dataset Mirror (for additional reference and exploration)***
  https://www.kaggle.com/datasets/uciml/default-of-credit-card-clients-dataset